

多体問題の計算科学

Computational Science for Many-Body Problems

#13 Parallelization for many-body problems
15:10-16:40 July 13, 2021

Lecture slide:

<https://github.com/compsci-alliance/many-body-problems>

1. Parallelization
2. Other numerical methods

Parallelization

(Revisited) Supercomputers in UTokyo

Ohtaka (CPU server)

ISSP, UTokyo@Kashiwa

Theoretical Peak: 6,881 TFlop/s

AMD EPYC 7720 2.0GHz, 64 cores x 2

Memory: 258GiB

1,680 nodes



東京大学 物性研究所
THE INSTITUTE FOR SOLID STATE PHYSICS
THE UNIVERSITY OF TOKYO

Oakforest-PACS

JCAHPC@Kashiwa (Utokyo & Tsukuba U)



Theoretical Peak: 24,913.5 TFlop/s

Intel Xeon Phi 7250 (1.4GHz, 68 cores, 112GB)

8,208 nodes

Oakbridge-CX (You can use this when you join CSA!)

Intel Xeon

1,368 nodes

(Revisited) “Fugaku” Supercomputer

Fugaku 富岳 @RIKEN, Kobe

-158,976 nodes

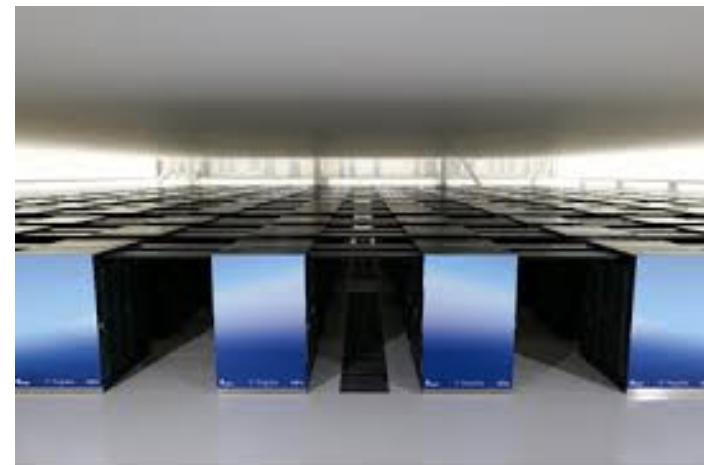
A64FX

48 cores/node

32 GiB/node

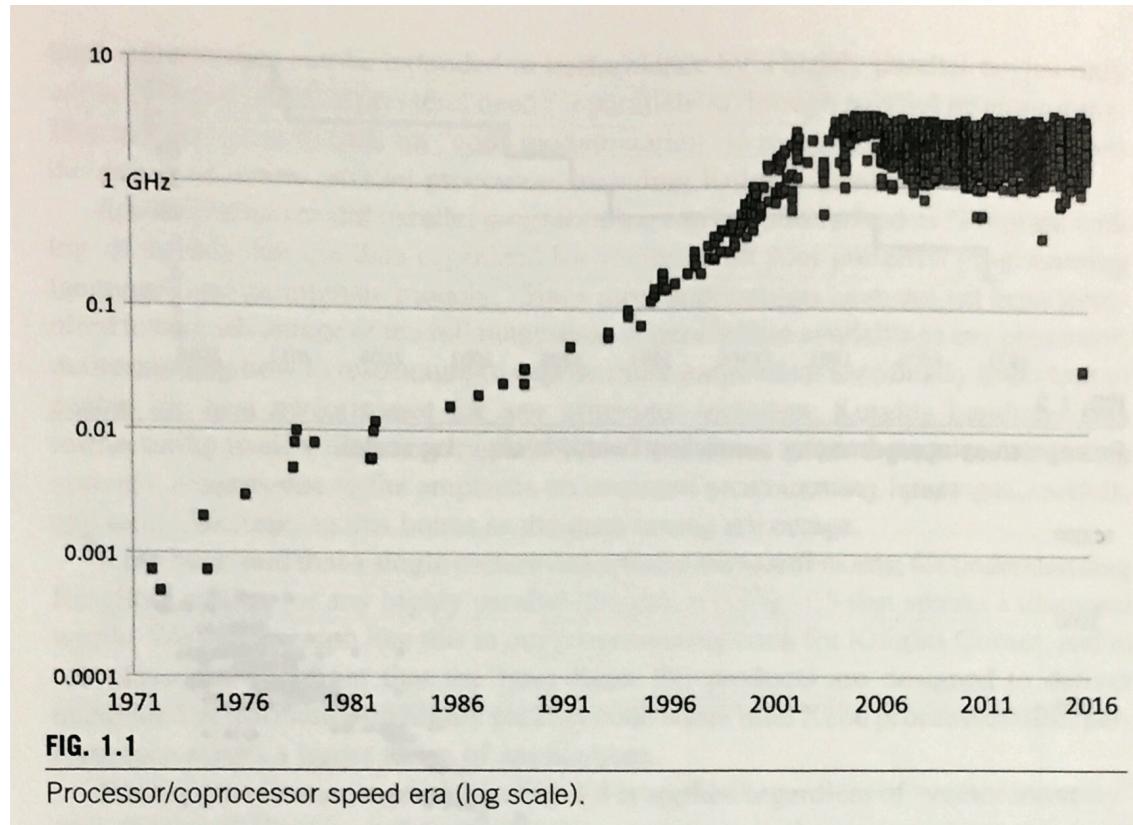


- 1st in Top500 & 9th in Green500 on June 2020
- 1st in Top500 & 10th in Green500 on November 2020



(Revisited) Increasing FLOP/s

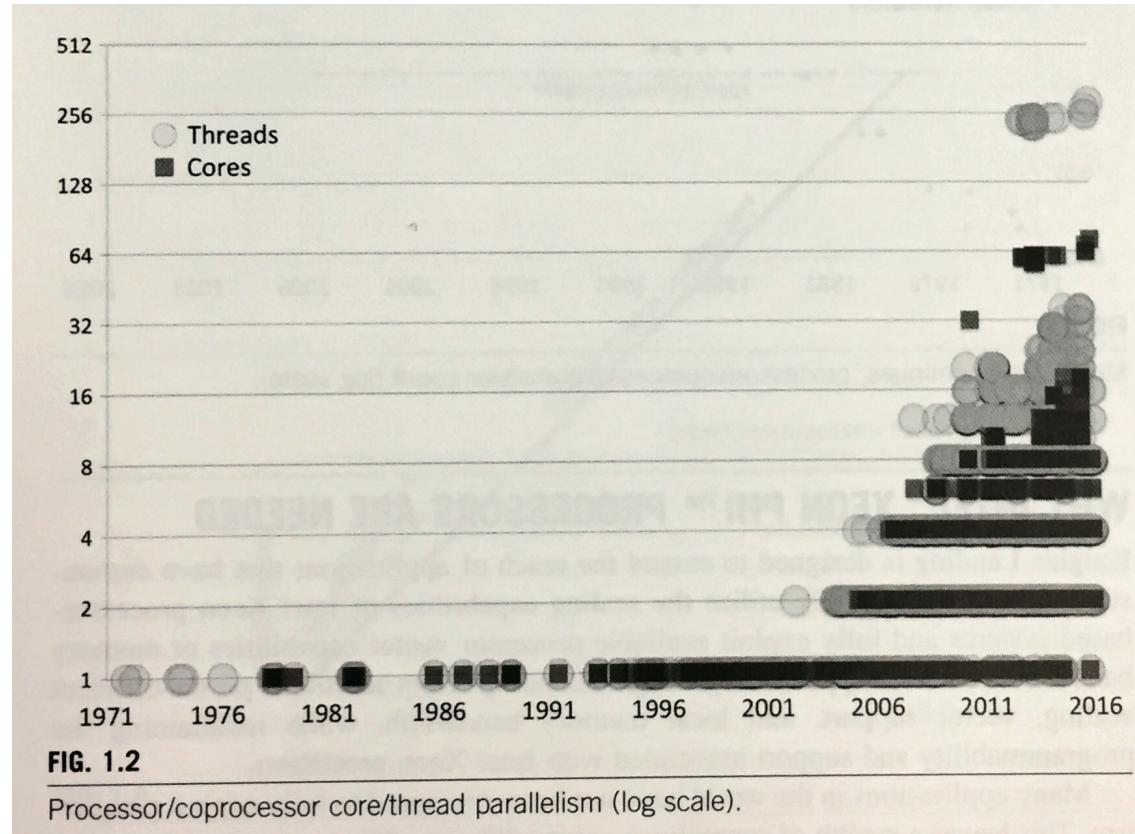
Clock rate saturated



J. Jeffers, J. Reinders, and A. Sodani,
Intel Xeon Phi Processor High Performance Programming

(Revisited) Increasing FLOP/s

Number of cores



J. Jeffers, J. Reinders, and A. Sodani,
Intel Xeon Phi Processor High Performance Programming

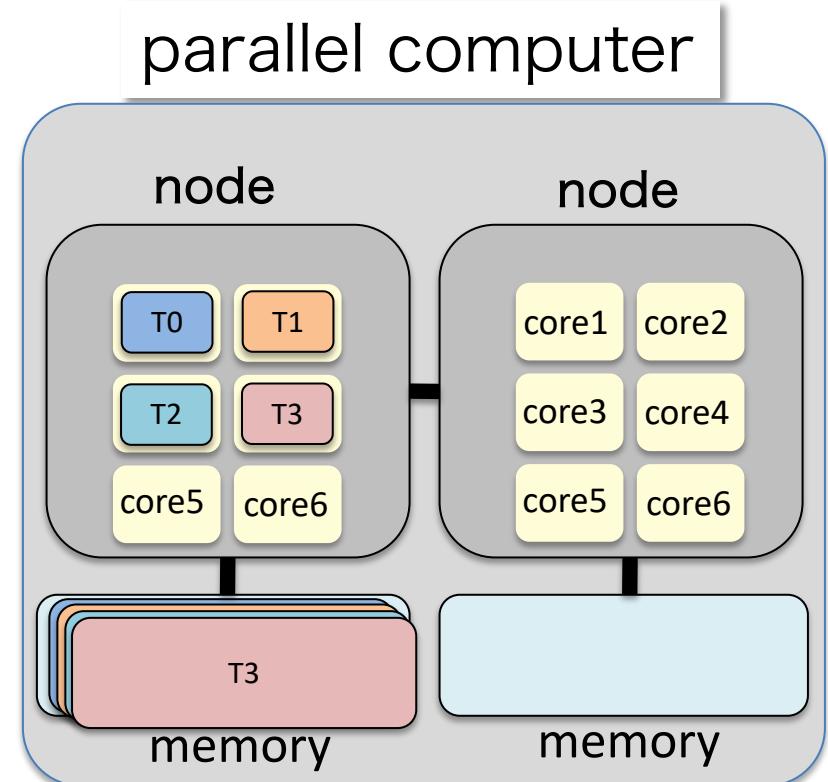
Parallelization

- Shared Memory Parallelization

Data on the memory of the node
is shared by the cores in the node

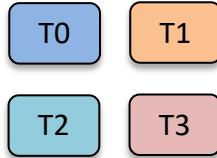
- Thread parallel

Realized by, for example,
OpenMP



Simple example in C lang: Divide a loop

```
#pragma omp parallel default(None) private(j) shared(a,b){  
#pragma omp for  
    for(j = 0; j < N; j++){  
        C[j] = a*A[j] + b;  
    }  
}
```



Dangerous example

```
#pragma omp parallel default(None) private(j) shared(a,b){  
#pragma omp for  
    for(j = 0; j < N-1; j++){  
        C[j+1] = a*C[j] + b;  
    }  
}
```

Dependency on former steps

Parallelization

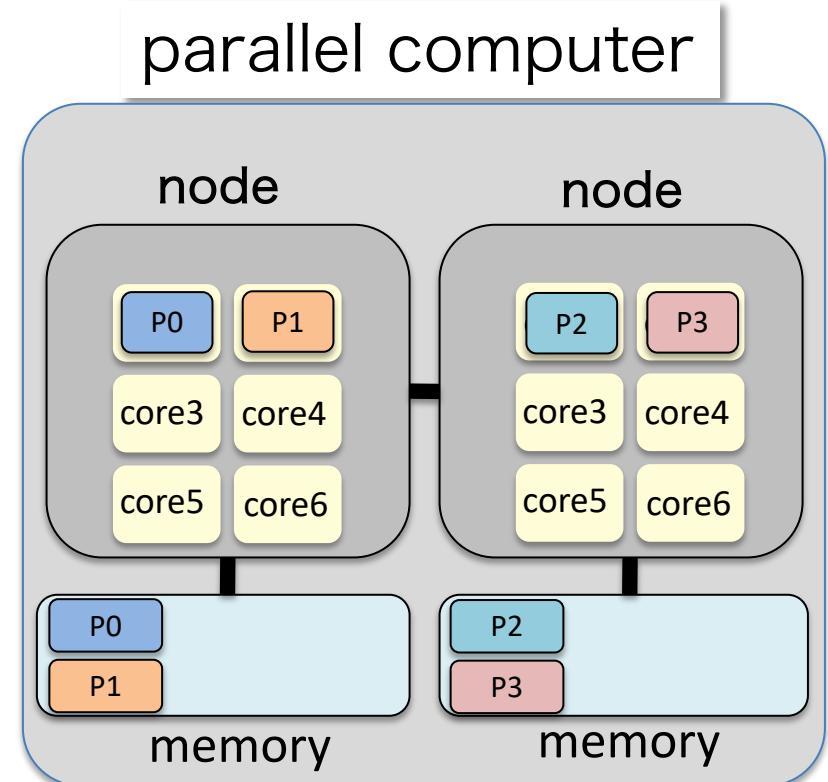
- Distributed Memory Parallelization

Each *process* can access its own data region in the node.

-Process parallel

Among processes,
data can be transferred
through MPI

(Message Passing Interface)



Declaration of MPI in C lang

```
ierr = MPI_Init(&argc, &argv);
ierr = MPI_Comm_size(MPI_COMM_WORLD, &nproc);
ierr = MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
```

nproc: # of processes

myrank: rank of the process, often used to generate
process-dependent random seed for MCMC

A Problem in Massively Parallel MCMC

Burn-in steps in Markov chain Monte Carlo

- Markov chain Monte Carlo (MCMC), for example, Metropolis-Hastings algorithm generates samples following a give probability distribution with *an initial state*
 - A finite sequence of steps in the initial stage of the Markov chain is *often* discarded:
Equilibration steps, burn-in or warm-up steps

cf.) Textbooks

A Problem in Massively Parallel MCMC

Extreme limit of massively parallel MCMC

- Few or few tens of warm-up steps and a single sample per process:
1 step annealed importance sampling with trivial prior distribution and no weights

of warm-up steps (> “auto correlation time”) is constant independent of # of processes

→ Total # of MC samples / # of processes will be much smaller than # of warm-up steps

Non Markov chain MC

Yukito Iba, *Population Monte Carlo algorithms*
<https://doi.org/10.1527/tjsai.16.279>

Population MC

- Diffusion MC
- Population annealing:
Annealed importance sampling+resampling

K. Hukushima and Y. Iba, AIP Conf. Proc. 690, 200 (2003)

Annealed importance sampling

$\{m_0^k, m_1^k, \dots, m_n^k\}$ D. A. Hendrix and C. Jarzynski, J. Chem. Phys. 114, 5974 (2001).
R. M. Neal, Statistics and Computing, 11, 125 (2001)

0. Prepare replicas sampled from a prior distribution with β_0
1. Change β_j ($j=0, 1, \dots$), calculate the weight, and anneal the replica and repeat 1.

$$W_j^k = W_{j-1}^k e^{-(\beta_j - \beta_{j-1}) f(m_{j-1}^k)}$$

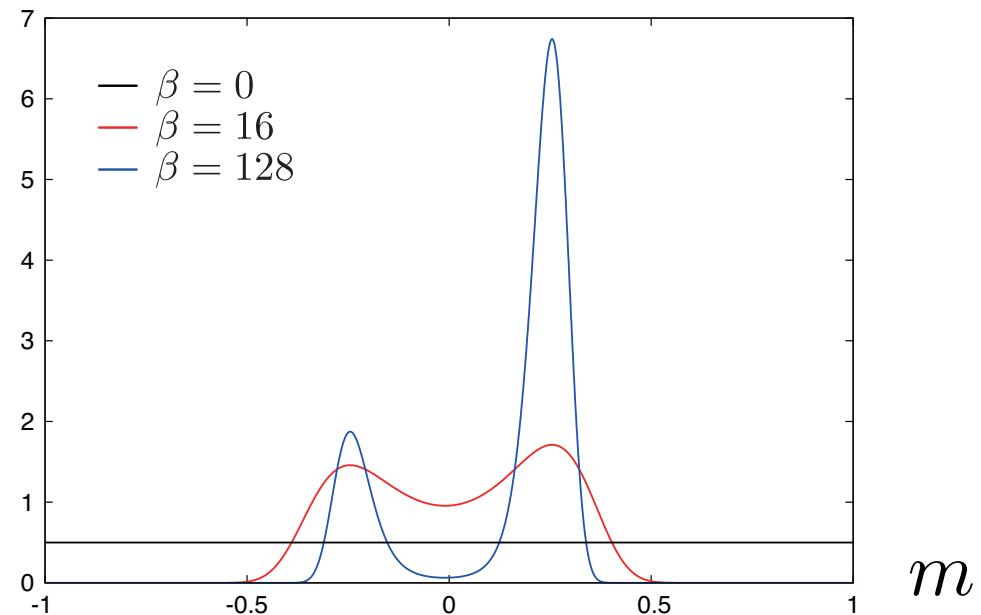
2. When $\beta_n = \beta_t$, take averages over the replicas with the weight

$$\beta_0 \leq \beta_1 \leq \dots \leq \beta_n = \beta_t$$

$$p(m) \propto e^{-\beta f(m)}$$

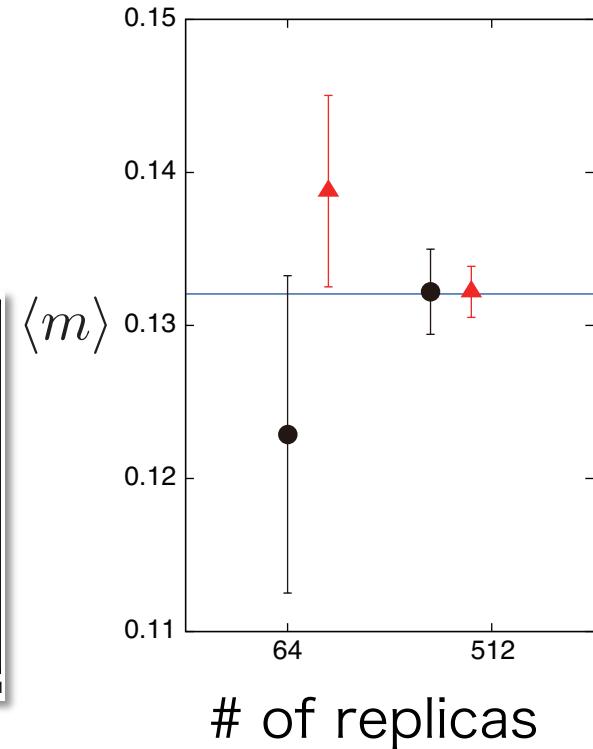
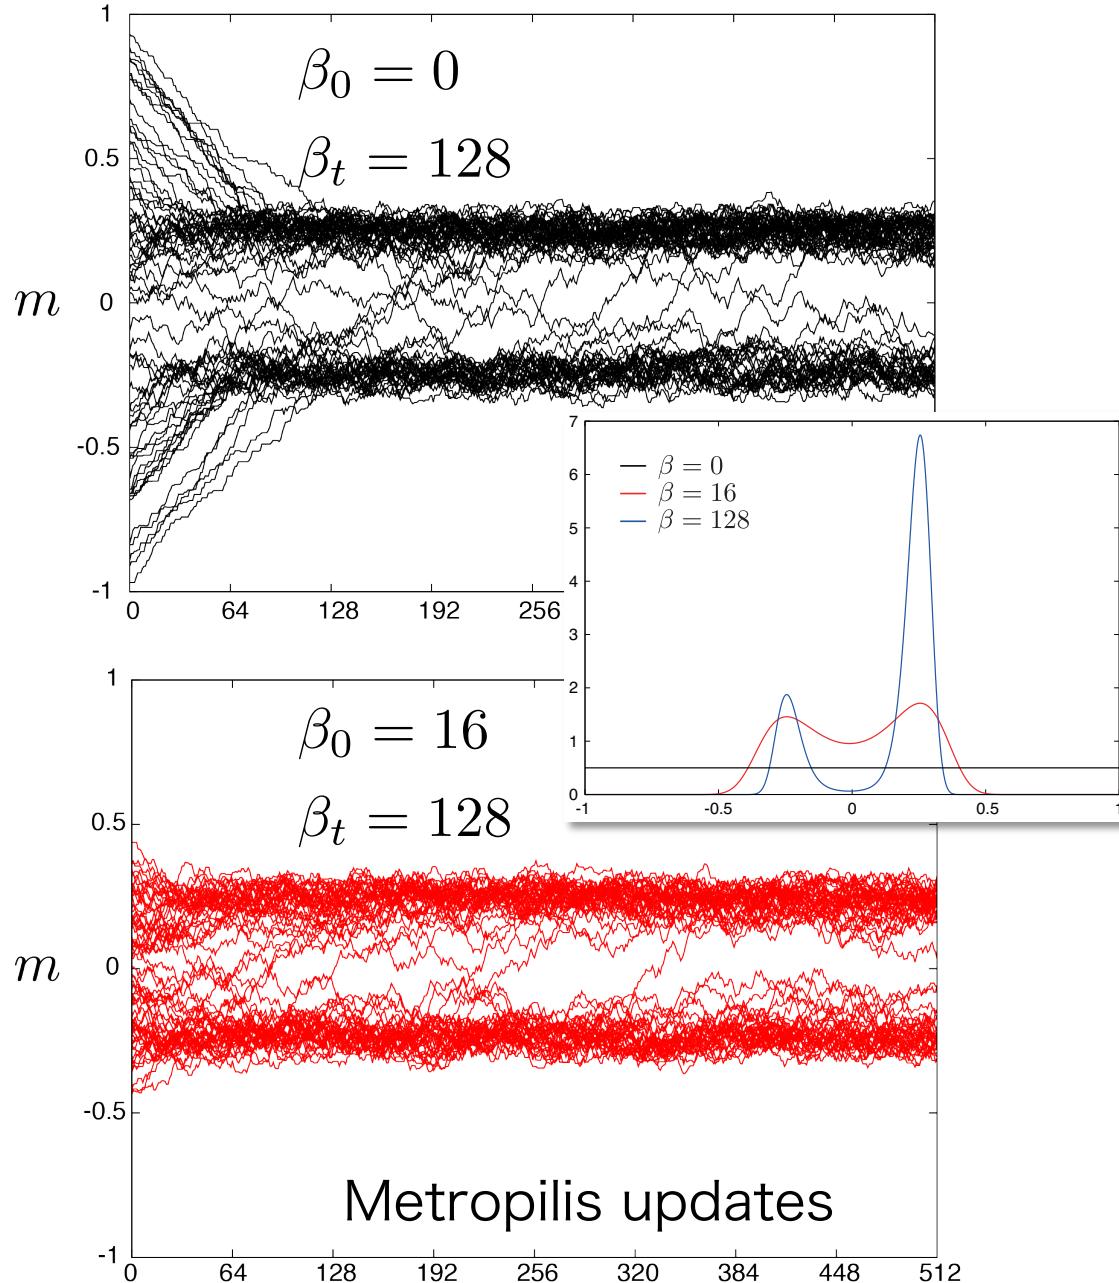
$$f(m) = -hm + rm^2 + vm^4$$

$$(h, r, v) = (0.02, -1, 8)$$



-Support condition $V_0 \supseteq V_1 \supseteq \dots \supseteq V_n$ $V_j = \{x | f_j(x) \neq 0\}$

One step annealed importance sampling



$$n = 1$$

$$W_1^k = e^{-(\beta_t - \beta_0)f(m_0^k)}$$

$$\langle m \rangle = \frac{\sum_k m_n^k W_n^k}{\sum_k W_n^k}$$

Distributed Wave Functions

Multiplication of Hamiltonian to Wave Function

Example of multiplication

-4 spins

$$(\hat{S}_1^- \hat{S}_3^+ + \hat{S}_1^+ \hat{S}_3^-) |0\textcolor{red}{1}00\rangle = |0\textcolor{red}{0}01\rangle$$

$$(\hat{S}_1^- \hat{S}_3^+ + \hat{S}_1^+ \hat{S}_3^-) |00\textcolor{red}{0}1\rangle = |01\textcolor{red}{0}0\rangle$$

$$\hat{S}_\ell^+ = S_\ell^x + iS_\ell^y$$

$$\hat{S}_\ell^- = S_\ell^x - iS_\ell^y$$

$$|\phi\rangle = \sum_{I_j \in \{0,1\}} C_{I_0 I_1 I_2 I_3} |I_0 I_1 I_2 I_3\rangle$$

$$(\hat{S}_1^- \hat{S}_3^+ + \hat{S}_1^+ \hat{S}_3^-) |\phi\rangle$$

$$C_{I_0 \textcolor{red}{1} I_2 0} \leftarrow C_{I_0 \textcolor{red}{0} I_2 1}$$

$$C_{I_0 0 I_2 1} \leftarrow C_{I_0 \textcolor{red}{1} I_2 0}$$

Multiplication of Hamiltonian to Wave Function

Example of multiplication

-4 spins

-Parallel: 2 processes

(processes are labeled by their rank)

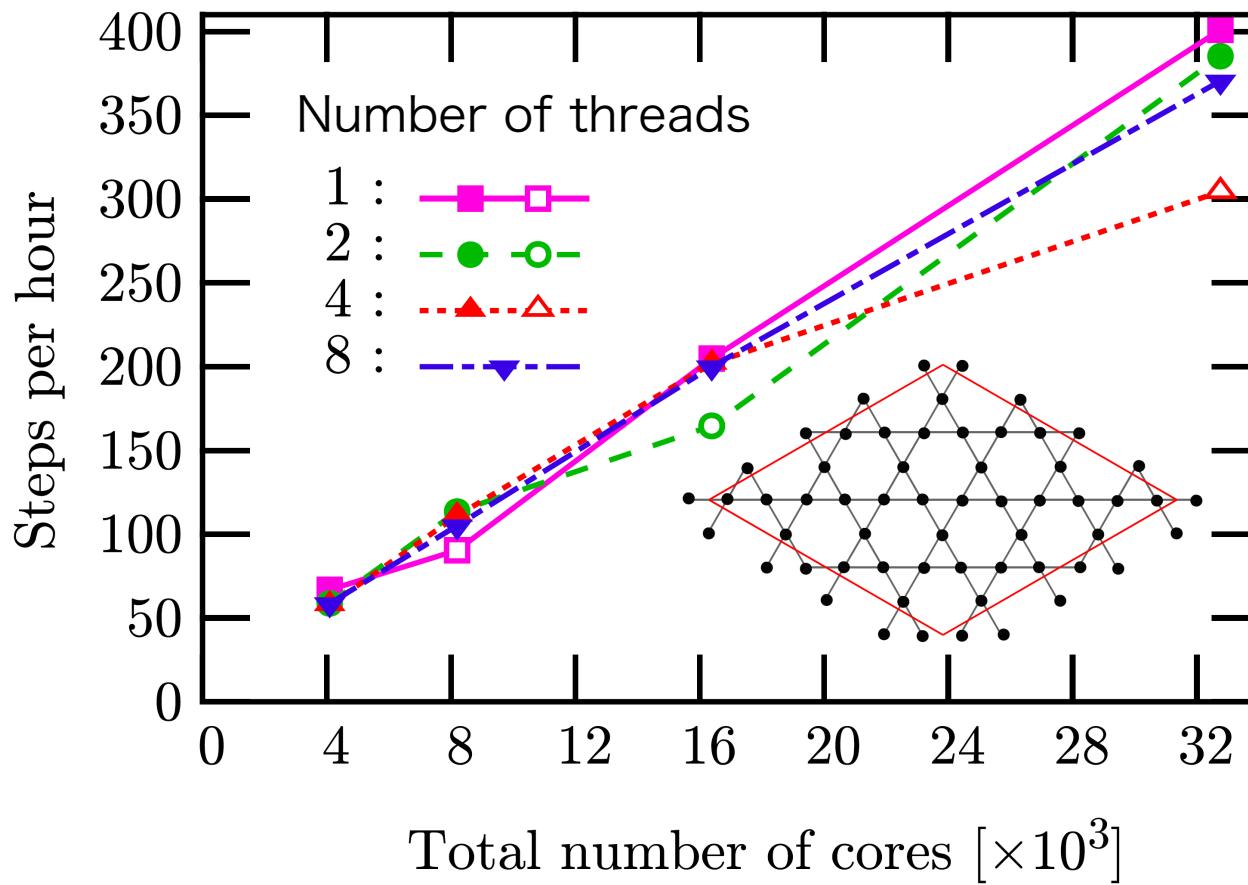
$$|I_0 I_1 I_2 \textcolor{blue}{I}_3\rangle$$

0	0000⟩	→ 0	
1	1000⟩	→ 1	
2	0100⟩	→ 8	
3	1100⟩	→ 9	
4	0010⟩	→ 4	rank 0
5	1010⟩	→ 5	
6	0110⟩	→ 12	
7	1110⟩	→ 13	
8	0001⟩	→ 2	
9	1001⟩	→ 3	
10	0101⟩	→ 10	
11	1101⟩	→ 11	
12	0011⟩	→ 6	rank 1
13	1011⟩	→ 7	
14	0111⟩	→ 14	
15	1111⟩	→ 15	

Speedup by Parallelization

36 spin Heisenberg model by $H\Phi$

@K computer, AICS, RIKEN



Other Numerical Methods

Other Numerical Methods

Ab initio

- Hartree-Fock theory*
- Density functional theory
 - Local density approximation*
 - Generalized gradient approximation*
 - Hybrid functional*
 - GW*
- Post Hartree-Fock (Quantum Chemistry)
 - Møller-Plesset*/Configuration interaction/Coupled cluster

*Difficulties in describing Mott insulators

- Transcorrelated

Other Numerical Methods

Many-body

- Exact diagonalization
- Quantum Monte Carlo (QMC) method
 - BSS, continuous time,...
 - Variational Monte Carlo
 - Green's function Monte Carlo
 - Diffusion Monte Carlo (*also used in *ab initio* approaches*)
- Numerical renormalization group (NRG) method
 - K. Wilson*
 - Density matrix renormalization group (DMRG) method
 - S. White*
 - Matrix product and tensor network method
 - *lecture in A term*
 - Dynamical mean-field theory

Combination

(Revisited) 2nd Report Problems
Please choose one of three

Report 2

Problem 1: Monte Carlo for quantum systems

1-1 (compulsory).

-Evaluate statistical errors in a 2-point distribution function $g(r)$ of liquid helium 4 at each distance r .

You may use McMillan's variational Jastrow wave function and variational parameters given in `vmc_helium4_2021.ipynb`.

-Obtain the relationship between the statistical errors and numbers of Monte Carlo samples (and confirm the error is proportional to $1/N_{MC}^{1/2}$, where N_{MC} is the number of Monte Carlo samples).

1-2 (optional).

-Obtain 2-point distribution functions $g(r)$ of liquid helium 3 for likewise (not likewise) spin pairs at ambient pressure.

You may use variational Jastrow-Slater wave functions and variational parameters obtained in D. Ceperley, G. V. Chester, and M. H. Kalos, Phys. Rev. B 16, 3081 (1977).

Report 2

Problem 2: Krylov subspace method

2-1 (compulsory).

-Implement CG and solve a linear question.

Please select a symmetric positive-definite matrix A with non-zero offdiagonal element.

The vector b should not be the zero vector.

Please choose A and b with N_H (linear dimension of A) = 5.

-Illustrate convergence of the solution obtained at each CG step.

You may plot step dependence of 2-norm of residual vectors r_k .

-Obtain the solution of $Ax=b$ by Lapack and compare with the solution by CG.

2-2 (Optional). Implement LOBCG method.

-Confirm that the code can calculate eigenvalues and eigenvectors for a real symmetric matrix.

-Solve an eigenvalue problem with $N_H > 5$

-Obtain the solution by Lapack and compare with the solution by CG.

Report 2

Problem 3: Open source software

3-1 (compulsory).

-Solve the following problems. You may use HΦ.

Estimate the energy difference between the lowest and 2nd lowest eigenstates of the 1dimensional $S=1/2$ and $S=1$ Hesenberg models with periodic boundary conditions.

-Use several L (number of spins) and extrapolate the gap to thermodynamic limit ($L \rightarrow \infty$). Please compare the extrapolated values with other results in the literature.

You may use $a+b/L+c/L^2$ or $a+b \exp(-cL)$.

-Illustrate the extrapolations.

3-2 (optional). Reproduce the figures in Sec. 2.2 of the tutorials of HΦ: [https://issp-center-dev.github.io/HPhi/
manual/develop/tutorial/en/html/finite_temperature/Kitaev.html](https://issp-center-dev.github.io/HPhi/manual/develop/tutorial/en/html/finite_temperature/Kitaev.html)

-Please choose one out of the three problems

-The code should be included (a jupyter notebook is recommended).

-Deadline: 8/6

Report 1 & 2

Deadline for Report 1:

-2020/7/13

Deadline for Report 2:

-2020/8/6

Please submit your report through ITC-LMS
(web page for the reports will be opened).

If you have any trouble, please contact us via email.

*If you are interested in QMC, open source software ALF & DSQSS would be worth trying.

Lecture Schedule

Classical

Quantum

- #1 Many-body problems in physics
- #2 Why many-body problem is hard to solve
- #3 Classical statistical model and numerical simulation
- #4 Classical Monte Carlo method and its applications
- #5 Molecular dynamics and its application
- #6 Extended ensemble method for Monte Carlo methods
- #7 Quantum lattice models and numerical approaches
- #8 Quantum Monte Carlo methods
- #9 Applications of quantum Monte Carlo methods
 - Path integral & applications
- #10 Linear algebra of large and sparse matrices for quantum many-body problems
- #11 Krylov subspace methods and their applications to quantum many-body problems
- #12 Large sparse matrices and quantum statistical mechanics
- #13 Parallelization for many-body problems